

CryptoStreamX:

Real-time Crypto Trade

Analytics Pipeline



Name: Nuwan Keshara Galappaththi

LinkedIn: <https://www.linkedin.com/in/nuwan-keshara/>

Contents

1) Project	3
2) Objectives / Problem Statement.....	3
3) Tech Stack.....	4
4) Architecture / System Design.....	5
Workflow Steps:	5
5) Implementation Steps	6
1. Data Ingestion (Python)	6
2. Data Quality Checks (Python)	6
3. Transformation and Modeling (ksqlDB)	6
4. Storage (ClickHouse).....	8
5. Code and Results	9
6) Summary	13

1) Project

- **CryptoStreamX: Real-time Crypto Trade Analytics Pipeline**

CryptoStreamX is a fully streaming data pipeline that captures and processes live cryptocurrency trade data from Binance in real time. Built with Apache Kafka, the system ingests, buffers and distributes high-throughput event streams. Kafka Streams and ksqlDB power in-flight transformations, aggregations and real-time analytics without external compute layers. Kafka Connect integrates with ClickHouse, a high-performance columnar database, for ultra-fast storage and query execution on time-series trade data. This architecture delivers a scalable, fault-tolerant data flow from ingestion to analytics, enabling instant insights into trading activity, price movements and market patterns.

2) Objectives / Problem Statement

- Live cryptocurrency trade data is highly volatile and comes as a continuous stream. Analyzing this data using traditional batch processes introduces latency, limiting real-time insights into trading activity and price fluctuations..
- Data Stream (Websocket) -
`wss://stream.binance.com:9443/stream?streams=btcusdt@trade`

- *Goal:*

Build a fully streaming data pipeline that:

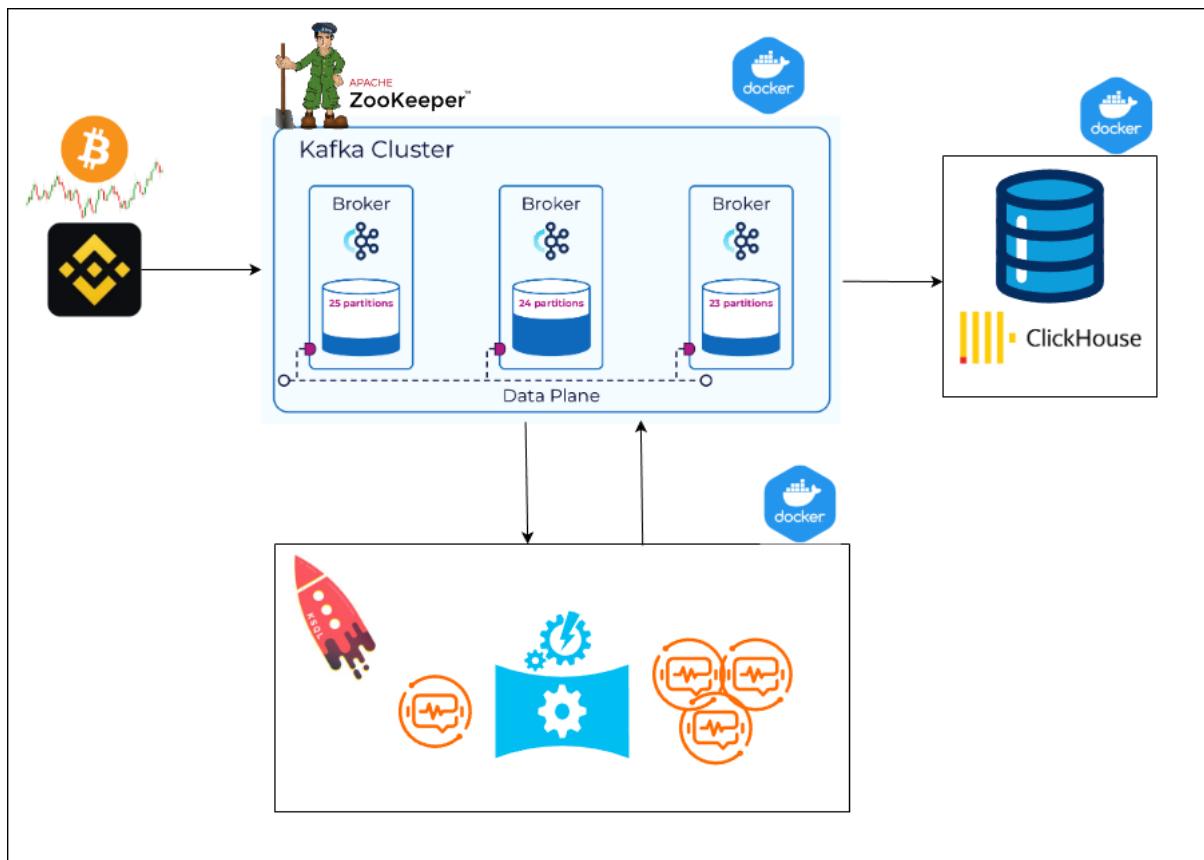
- Ingests live trade data from Binance in real time.
- Ensures data quality and schema consistency via Avro schemas and Kafka.
- Applies in-flight transformations and aggregations with ksqlDB.
- Stores processed trade data in ClickHouse for high-performance analytics.
- Enables near-real-time reporting.

3) Tech Stack

Layer	Tool	Purpose
Data Source	Binance WebSocket API	Live cryptocurrency trades (BTC, ETH, etc.)
Messaging & Streaming	Apache Kafka	High-throughput message broker for real-time ingestion and buffering
Cluster Management	Apache Zookeeper	Manage Cluster brokers
Schema Management	Confluent Schema Registry	Enforces Avro schemas for data consistency
Streaming Processing	ksqlDB / Kafka Streams	In-flight data transformations, aggregations and windowed analytics
Storage / Analytics	ClickHouse	Columnar DB optimized for time-series and high-performance queries
Programming	Python (WebSockets, Confluent Kafka)	Data ingestion, transformations
Version Control	Git / GitHub	Code management
Environment Management	Python dotenv	Secure configuration management

4) Architecture / System Design

- Pipeline Flow



Workflow Steps:

- I. Live ingestion from Binance WebSocket into Kafka topics.
- II. Schema validation using Confluent Schema Registry to ensure Avro schema compliance.
- III. In-flight transformations and aggregations in ksqlDB:
- IV. Windowed aggregations (1-minute and 5- minute tumbling windows, 10-second grace periods)
- V. Real-time calculation of metrics like trade count, total quantity, average price, min/max prices.
- VI. High-performance storage in ClickHouse for real-time querying.

5) Implementation Steps

1. Data Ingestion (Python)

- WebSocket Client:
 - Connects to Binance WebSocket API.
 - Receives live trade messages and parses JSON payloads.
 - Transforms each trade into Avro-compliant dictionary records.
- Kafka Producer:
 - Serializes messages with AvroSerializer.
 - Publishes trade data to Kafka topic ‘binance_trade’.
 - Handles errors, retries and ensures minimal message loss.

2. Data Quality Checks (Python)

- Validates messages for:
 - Missing fields (e.g., event_time, symbol, trade_id)
 - Numeric validation (price ≥ 0 , quantity ≥ 0)
 - Boolean field normalization (market_maker, ignore)
 - Logs issues to logs/producer.log.

3. Transformation and Modeling (ksqlDB)

- Stream Creation:

```
✓ CREATE STREAM BINANCE_TRADE_RAW (
    "event_type" VARCHAR,
    "event_time" BIGINT,
    "symbol" VARCHAR,
    "trade_id" BIGINT,
    "price" DOUBLE,
    "quantity" DOUBLE,
    "trade_time" BIGINT,
    "market_maker" BOOLEAN,
    "ignore" BOOLEAN
) WITH (
    KAFKA_TOPIC='binance_trade',
    VALUE_FORMAT='AVRO',
    TIMESTAMP='event_time'
);
```

- Aggregations (1-minute tumbling windows):

```

CREATE TABLE BINANCE_TRADE_AGG_1M
WITH (
    KAFKA_TOPIC='binance_trade_agg_1m',
    VALUE_FORMAT='AVRO'
) AS
SELECT
    "symbol" AS symbol,
    CAST(WINDOWSTART AS BIGINT) AS "window_start",
    CAST(WINDOWEND AS BIGINT) AS "window_end",
    CAST(COUNT(*) AS BIGINT) AS "trade_count",
    CAST(SUM("quantity") AS DOUBLE) AS "total_quantity",
    CAST(AVG("price") AS DOUBLE) AS "avg_price",
    CAST(MIN("price") AS DOUBLE) AS "min_price",
    CAST(MAX("price") AS DOUBLE) AS "max_price"
FROM BINANCE_TRADE_RAW
WINDOW TUMBLING (SIZE 1 MINUTE, GRACE PERIOD 10 SECONDS)
GROUP BY "symbol"
EMIT FINAL;

```

- Aggregations (5-minute tumbling windows):

```

CREATE TABLE BINANCE_TRADE_AGG_5M
WITH (
    KAFKA_TOPIC='binance_trade_agg_5m',
    VALUE_FORMAT='AVRO'
) AS
SELECT
    "symbol" AS symbol,
    CAST(WINDOWSTART AS BIGINT) AS "window_start",
    CAST(WINDOWEND AS BIGINT) AS "window_end",
    CAST(COUNT(*) AS BIGINT) AS "trade_count",
    CAST(SUM("quantity") AS DOUBLE) AS "total_quantity",
    CAST(AVG("price") AS DOUBLE) AS "avg_price",
    CAST(MIN("price") AS DOUBLE) AS "min_price",
    CAST(MAX("price") AS DOUBLE) AS "max_price"
FROM BINANCE_TRADE_RAW
WINDOW TUMBLING (SIZE 5 MINUTE, GRACE PERIOD 10 SECONDS)
GROUP BY "symbol"
EMIT FINAL;

```

4. Storage (ClickHouse)

- Kafka Connect Sink streams aggregated tables into ClickHouse.
 - ClickHouse schema:

```
CREATE DATABASE IF NOT EXISTS crypto;
USE crypto;

CREATE TABLE binance_trade (
    event_type      String,
    event_time     UInt64,
    trade_time     UInt64,
    symbol         String,
    trade_id       UInt64,
    price          Float64,
    quantity       Float64,
    market_maker   Bool,
    ignore         Bool,
    event_time_dt  DateTime64(3, 'UTC') ALIAS toDateTime(event_time / 1000, 'UTC'),
    trade_time_dt  DateTime64(3, 'UTC') ALIAS toDateTime(trade_time / 1000, 'UTC')
)
ENGINE = MergeTree()
ORDER BY (event_time);
```

- Aggregations (1-minute tumbling windows):

```
CREATE TABLE binance_trade_agg_1m (
    window_start    UInt64,
    window_end     UInt64,
    trade_count    UInt64,
    total_quantity Float64,
    avg_price      Float64,
    min_price      Float64,
    max_price      Float64,
    window_start_dt DateTime64(3, 'UTC') ALIAS toDateTime(window_start / 1000, 'UTC'),
    window_end_dt  DateTime64(3, 'UTC') ALIAS toDateTime(window_end / 1000, 'UTC')
)
ENGINE = MergeTree()
ORDER BY (window_start);
```

- Aggregations (5-minute tumbling windows):

```
CREATE TABLE binance_trade_agg_5m (
    window_start    UInt64,
    window_end     UInt64,
    trade_count    UInt64,
    total_quantity Float64,
    avg_price      Float64,
    min_price      Float64,
    max_price      Float64,
    window_start_dt DateTime64(3, 'UTC') ALIAS toDateTime(window_start / 1000, 'UTC'),
    window_end_dt  DateTime64(3, 'UTC') ALIAS toDateTime(window_end / 1000, 'UTC')
)
ENGINE = MergeTree()
ORDER BY (window_start);
```

5. Code and Results

- Producer Example:

```
# Kafka Producer
producer_conf = {
    'bootstrap.servers': BOOTSTRAP_SERVERS,
    'key.serializer': StringSerializer('utf_8'),
    'value.serializer': avro_serializer,
    'compression.type': 'lz4',
    'linger.ms': 20,
    'batch.size': 32768,
    'acks': '1'
}

# Message formatter
def make_record(payload):
    """Convert Binance trade message to Avro record"""
    data = payload.get("data", payload)
    if not data:
        return None

    try:
        return {
            "event_type": data.get("e"),
            "event_time": int(data.get("E", 0)),
            "symbol": data.get("s"),
            "trade_id": int(data.get("t", 0)),
            "price": float(data.get("p", 0.0)),
            "quantity": float(data.get("q", 0.0)),
            "trade_time": int(data.get("T", 0)),
            "market_maker": bool(data.get("m", False)),
            "ignore": bool(data.get("M", False))
        }
    except Exception as e:
        logging.error(f"Failed to parse record: {e}")
        return None
```

- Kafka UI:

The screenshot shows the Kafka UI Dashboard. On the left, there's a sidebar with links for 'Dashboard', 'cryptostream...', 'Brokers', 'Topics', and 'Consumers'. The main area is titled 'Dashboard' and shows the following cluster statistics for 'cryptostream-kafka-cluster':

- Clusters: 1 online, 0 offline
- Version: 3.9-IV0
- Brokers count: 3
- Partitions: 142
- Topics: 16
- Production: 0 Bytes
- Consumption: 0 Bytes

Buttons at the bottom right include 'Configure new cluster' and 'Configure'.

The screenshot shows the Kafka UI Brokers page. It displays broker information and partition skew across three brokers:

Uptime			Partitions			
Broker Count	Active Controller	Version	Online	URP	In Sync Replicas	Out Of Sync Replicas
3	2	3.9-IV0	142 of 142	0	249 of 249	0
			Broker ID	Disk usage	Partitions skew	Leaders
			0	1.34 MB, 84 segment(s)	1.20%	48
			1	340.01 KB, 81 segment(s)	-2.40%	46
			2	291.1 KB, 84 segment(s)	1.20%	48

- Kafka Topics:

The screenshot shows the Kafka UI Topics page. It lists the following topics:

Topic Name	Partitions	Out of sync replicas
binance_trade	1	0
binance_trade_agg_1m	1	0
binance_trade_agg_5m	1	0
crypto_ksqldbksql_processing_log	1	0
dlq-clickhouse	1	0

At the top, there are buttons for 'Delete selected topics', 'Copy selected topic', and 'Purge messages of selected topics'. There's also a search bar and a 'Show Internal Topics' toggle switch.

- Raw data stream- binance_trade topic

Topics / binance_trade

Overview	Messages	Consumers	Settings	Statistics
Seek Type	Offset	Partitions	Key Serde	Value Serde
Offset	Offset	All items are selected.	String	String
Search <input type="text"/>		+ Add Filters	Clear all <input type="button" value="Submit"/>	
Offset	Partition	Timestamp	Key	Preview
0	0	11/7/2025, 14:49:29	BTCUSDT	
1	0	11/7/2025, 14:49:30	BTCUSDT	
2	0	11/7/2025, 14:49:30	BTCUSDT	
3	0	11/7/2025, 14:49:30	BTCUSDT	
4	0	11/7/2025, 14:49:30	BTCUSDT	
5	0	11/7/2025, 14:49:30	BTCUSDT	
6	0	11/7/2025, 14:49:30	BTCUSDT	
7	0	11/7/2025, 14:49:30	BTCUSDT	
8	0	11/7/2025, 14:49:30	BTCUSDT	

- Aggregated Data Stream 1-minute period - binance_trade_agg_1m (Avro)

Topics / binance_trade_agg_1m

Overview Messages Consumers Settings Statistics

Seek Type: Offset | Partitions: All items are selected. | Key Serde: String | Value Serde: String | Clear all | Submit

Search: + Add Filters

	Offset	Partition	Timestamp	Key	Preview
1	0	0	11/7/2025, 14:49:59	BTTCUSDT	███████^███
2	1	0	11/7/2025, 14:50:59	BTTCUSDT	███████]███
3	2	0	11/7/2025, 14:51:59	BTTCUSDT	███████^███
4	3	0	11/7/2025, 14:52:59	BTTCUSDT	███████^███
5	4	0	11/7/2025, 14:53:59	BTTCUSDT	███████^███

← Back Next →

- Aggregated Data Stream 5-minute period - binance_trade_agg_5m (Avro)

Offset	Partition	Timestamp	Key	Preview
0	0	11/7/2025, 14:49:59	BTCUSDT	[REDACTED]

- Clickhouse Database raw data stream

	AZ event_type	123 event_time	123 trade_time	AZ symbol	123 trade_id	123 price	123 quantity	market_maker	ignore
1	trade	1,762,520,465,734	1,762,520,465,734	BTCUSDT	5,451,030,795	99,621.41	0.00027	[v]	[v]
2	trade	1,762,520,465,700	1,762,520,465,699	BTCUSDT	5,451,030,794	99,621.42	0.00005	[]	[v]
3	trade	1,762,520,465,899	1,762,520,465,899	BTCUSDT	5,451,030,796	99,621.42	0.00039	[]	[v]
4	trade	1,762,520,465,242	1,762,520,465,241	BTCUSDT	5,451,030,772	99,621.42	0.00102	[]	[v]
5	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,773	99,621.41	0.00006	[v]	[v]
6	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,774	99,621.41	0.00006	[v]	[v]
7	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,775	99,621.41	0.00006	[v]	[v]
8	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,776	99,621.41	0.00006	[v]	[v]
9	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,777	99,621.41	0.00006	[v]	[v]
10	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,778	99,621.41	0.00006	[v]	[v]
11	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,779	99,621.41	0.00006	[v]	[v]
12	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,780	99,621.41	0.00006	[v]	[v]
13	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,781	99,621.41	0.00049	[v]	[v]
14	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,782	99,621.41	0.00077	[v]	[v]
15	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,783	99,621.41	0.00006	[v]	[v]
16	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,784	99,621.41	0.00191	[v]	[v]
17	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,785	99,621.41	0.00006	[v]	[v]
18	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,786	99,621.41	0.00006	[v]	[v]
19	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,787	99,621.41	0.00008	[v]	[v]
20	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,788	99,621.41	0.0167	[v]	[v]
21	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,789	99,621.41	0.00006	[v]	[v]
22	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,790	99,621.41	0.00007	[v]	[v]
23	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,791	99,621.41	0.00006	[v]	[v]
24	trade	1,762,520,465,268	1,762,520,465,267	BTCUSDT	5,451,030,792	99,621.41	0.00006	[v]	[v]

- Clickhouse aggregated data (1m and 5m)

	123 window_start	123 window_end	123 trade_count	123 total_quantity	123 avg_price	123 min_price	123 max_price
0	1,762,520,400,000	1,762,520,460,000	3,484	13.68888	99,617.6062399552	99,593.3	99,648.26
1	1,762,520,280,000	1,762,520,340,000	3,998	19.84713	99,654.618090492	99,614.82	99,689.87
2	1,762,519,740,000	1,762,519,800,000	3,379	13.72041	99,391.3232050896	99,366	99,423.91
3	1,762,519,800,000	1,762,519,860,000	7,058	28.29486	99,482.8056545781	99,411.76	99,550.11
4	1,762,519,860,000	1,762,519,920,000	5,816	19.46632	99,484.6170512392	99,434.86	99,547.21
5	1,762,519,920,000	1,762,519,980,000	8,550	28.79305	99,523.794521637	99,455.99	99,573.83
6	1,762,519,980,000	1,762,520,040,000	5,009	14.46967	99,485.1600399288	99,452.51	99,509.29
7	1,762,520,040,000	1,762,520,100,000	5,430	28.53431	99,408.8566740333	99,353.05	99,469.39
8	1,762,520,100,000	1,762,520,160,000	9,957	53.78408	99,368.6001837906	99,260.86	99,462.55
9	1,762,520,160,000	1,762,520,220,000	5,204	65.46886	99,533.562142583	99,452.58	99,606.8
10	1,762,520,220,000	1,762,520,280,000	5,090	30.07635	99,664.3543929255	99,590.72	99,704.49
11	1,762,520,340,000	1,762,520,400,000	5,241	40.06936	99,628.0618526995	99,580.2	99,668.89

6) Summary

CryptoStreamX demonstrates a modern, end-to-end streaming data architecture that is scalable, fault-tolerant and built for real-time analytics. This system captures live cryptocurrency trade data, validates and transforms it in-flight and delivers analytics-ready metrics with minimal latency.

Core Components:

- Apache Kafka:
 - Handles high-throughput, real-time data ingestion and buffering between producers and consumers.
- Apache Zookeeper:
 - Manages Kafka cluster coordination, leader election and configuration synchronization across brokers, ensuring high availability and fault tolerance.
- Confluent Schema Registry:
 - Maintains Avro schemas to enforce consistent message structure and backward compatibility between producer and consumer applications.
- ksqlDB:
 - Performs continuous, in-flight transformations and real-time aggregations directly on Kafka topics.
- ClickHouse:
 - A high-performance columnar database for storing and querying time-series trade data, enabling fast analytical queries and historical trend analysis.

Key Achievements:

- Real-time trade analytics with low latency from data ingestion to visualization.
- Fault-tolerant Kafka cluster managed by Zookeeper for reliable message delivery and partition replication.
- Schema validation via Avro, preventing ingestion of inconsistent messages.

- Fully modular architecture, easily extendable to new data sources or additional metrics.
- Instant analytical queries, supporting aggregations on trade data.